

CS180/280A Discussion #1

Konpat

Credits:

Justin, Chung Min

Welcome!!

GSI's



Justin Kerr



Konpat Preechakul



Chung Min Kim



Brent Yi

Tutors



Jameson Crate




Jorge Diaz Chao



Natalie Wei



Jingfeng Yang



Me: **Konpat** Preechakul 🖐️


“Learning abstractions from pixels”

Scene understanding



Visual Jenga

Diffusion models & Representation learning



Machine learning

Cellular automata



Physics simulations



Reminders

Proj1 due Thurs **9/12** 11:59pm

OH dates are released!

Worksheets online:

Topics

Discussion	Topic	Materials
Week 1	Python & NumPy Fundamentals for Computer Vision	Main sheet Challenge Solutions

Discussions this year!

- **Practical practice** (for Projs) + **Conceptual understanding** (for exams)
- **Collaborative!** Move to be near someone! :)
- **Minimal laptop.** We want you to go through with your hand!
 - For future sessions if you want to use laptop, please sit in the back.
- Numbered problems are in scope for exams, bonus questions are intended to be *hard* for people who want to try them
- **Note:** these are new this year! (rough ☐)
- After next week will circulate a feedback form

What are your questions?

Agenda

- Short lectures (10 mins)
- Problems (10 mins)
- A bit more lectures (10 mins)
- A bit more problems (10 mins)

This discussion: Visual world 🧡 Computers

“Intro to Computer Vision and Computational Photography”



???

This discussion: Visual world 🧡 Computers

We need to learn how to **1) input, 2) store, and 3) manipulate 4) output images!**



Storing data?



Visualizing?

What data structure are images?

An image is an **array** of pixels!

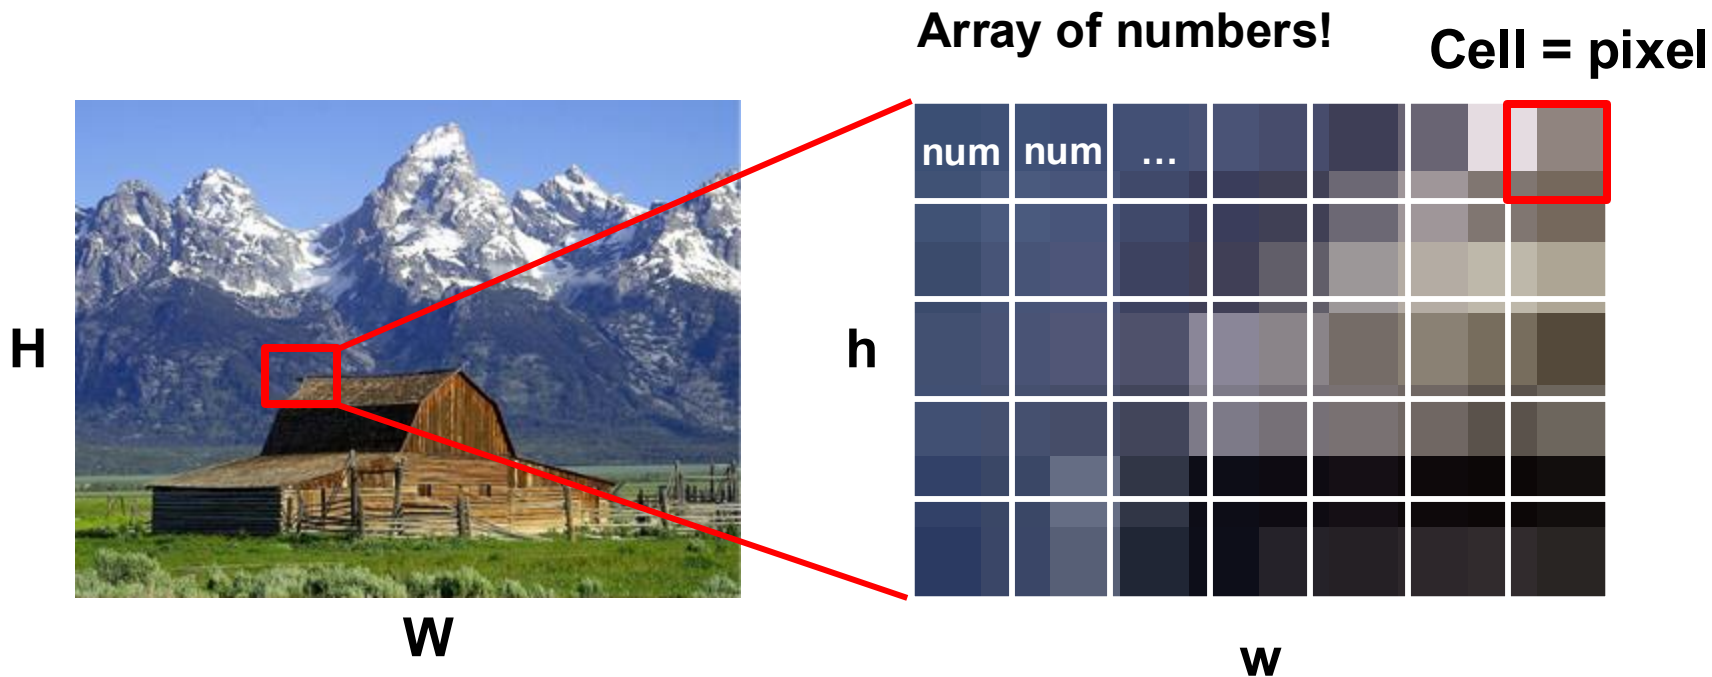


H


W

What data structure are images?

An image is an **array** of pixels!




What is a pixel?



*Range (0 – 255)
Each color 8 bits
Total **24-bit color**

COLOR!

3 channels: red, green, blue





*RGB vs. BGR conventions

Pixel is not always 3 numbers!


Grayscale!

1 channels: intensity



Pixel can be something else!


180 190 170



Depth!

3 (feet)

2.5 (feet)



How do we get these things into python anyway...

Notebook demo

we load with **cv2.imread(filename)**

we visualize with matplotlib's **plt.imshow(filename)**

then we probe like, what is the type / shape of this image

lol it is **an array!**

what are the values, etc.

zoom in, zoom out, get the colors right

Inspecting images

```
>>> img = cv2.imread("img.jpg")

>>> print(img.shape) # shape (1080, 1920, 3)

>>> print(img.dtype) # np.uint8

>>> print(img.min(), img.max()) # 0 255
```

2 primary formats:


- uint8, 0->255 scaling
- float, 0->1.0 scaling

Be careful converting!


Pixel layout in an array

```
>>> img = cv2.imread("img.jpg") # shape (1080, 1920, 3)
```

Pixels stored along **channels**.



(H, W, C) $C=3$ "channel-last"
... or (C, H, W) "channel-first"



(H, W) or $(H, W, 1)$

Singleton dim.

*BGR is a bit hard to visualize actually...

Good news: many image operations are just array operations!

NumPy



So fast, so easy 😊

Let's start: Color channel manipulation (1.3 Slicing)

```
>>> img
```


Original Image



150
101
105

```
>>> img[...,0]
```

Blue Channel



105

```
>>> img[...,1]
```

Green Channel



101

```
>>> img[...,2]
```

Red Channel




150

*Reminds of Proj 1!

BGR => RGB (1.3 Slicing)

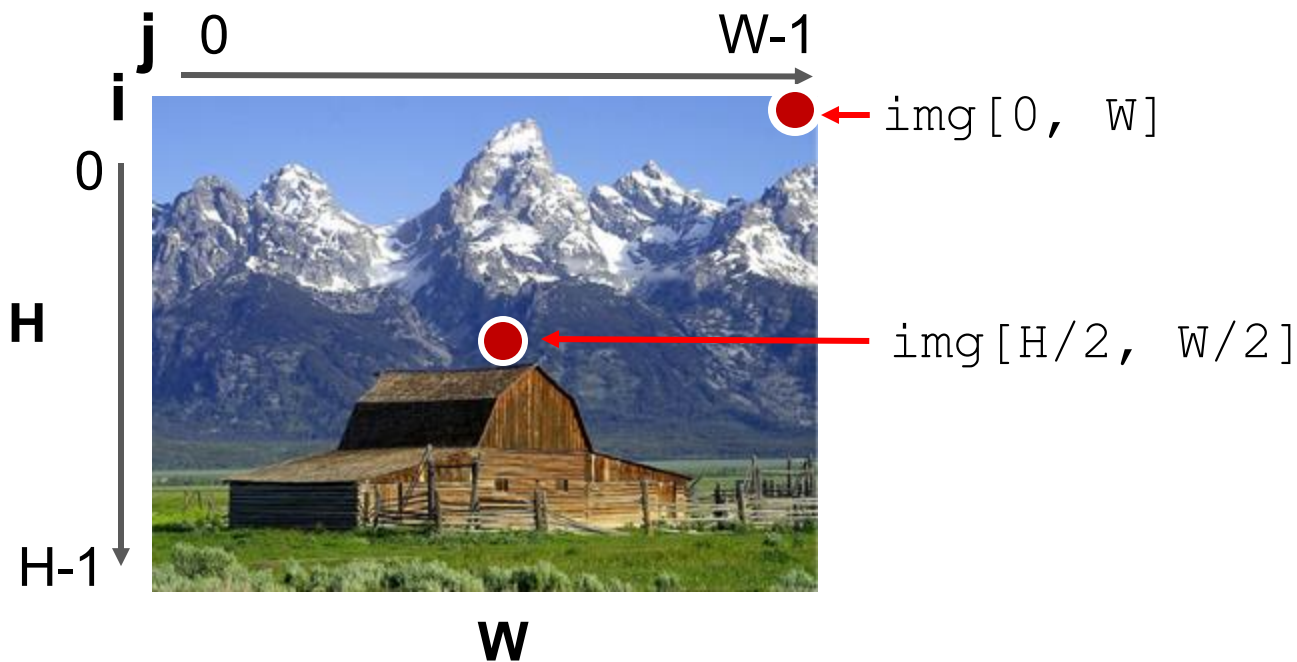
```
>>> img = img[:, :, [2,1,0]]
```



*Easier to plot. Show on notebook

Indexing conventions (1.3 Slicing)

Index into arrays like i, j in a matrix, **not like** x, y in a coordinate plane!



: (colon)


Cropping images (1.3 Slicing)

```
>>> left_half = img[:, :100, :]
```

```
>>> bottom_half = img[100:]
```

:100

100:



*Let me show you guys

Joining images (1.2 Stack & Concat)

```
>>> vertical = np.concatenate([angjoo,alyosha],axis=0)
```



angjoo



alyosha



Joining images (1.2 Stack & Concat)

```
np.concatenate([angjoo,alyosha],axis=0)
```




angjoo



alyosha

axis 1

axis 0




axis 0

```
np.concatenate([angjoo,alyosha],axis=1)
```

axis 1

axis 1

axis 0



*Let me show you guys.

What are videos? (1.2 Stack & Concat)

```
>>> video = np.stack(frames, axis=0)
```

Videos are just arrays of *batches* of images!



video
(T, H, W, C)



video [0]



video [5]



video [11]



video [15]

(H, W, C)

What are videos? (1.2 Stack & Concat)

```
>>> video = np.stack(frames, axis=0)
```

Videos are just arrays of *batches* of images!



**Video being
played**

(T, H, W, C)



Actual video

"Cube"

(T, H, W, C)

NumPy basics: do **Problems 1.1-1.8 (5 mins)** with the people around you!

Quick refresher on the following:

- `np.array([1, 2, 3])`
- `np.full(shape , value)`
- `array.astype(type)`
- **type:** `np.uint8`, `np.float32`, `np.float64`
- `array[i, j]`, `array[a : b]`
- `np.concatenate([a ,b], axis=?)`
- `np.stack([a, b], axis=?)`

(Then we will go over quickly)

Pixel operations: Do **Problems 2.1-2.5 & 2.9 (5 mins)**

(Then we will go over quickly).

What happens when shapes don't match? (3)



```
>>> img = cv2.imread("img.jpg") # shape (1080, 1920, 3)
```

```
>>> brighter_img = img + np.array([100,100,100]) # shape (1080,1920,3) + (3,) ??
```




*Let me show you.

*Overflow

Broadcasting: automatically *repeat* elements to match!

```
>>> a = np.array([1, 2, 3])      # shape (3,)  
>>> b = np.array(2)             # shape ()!  
>>> print( a * b )               # [2.0,4.0,6.0] shape (3,)
```



Broadcasting

Rule for figuring out behavior:

1. **Line up** array shapes starting from the **right**
2. **For each axis:**
 - a. If shapes match, continue to the left
 - b. If shapes don't match and one is 1, stretch its values to fit the larger
 - c. If shapes don't match and *neither* are 1, throw an error

$A = (2, 3)$
 $B = (2, 1, 1)$ \rightarrow $A = (2, 3)$
 $B = (2, 1, 3)$ \rightarrow $A = (2, 3)$
 $B = (2, 2, 3)$ \rightarrow $A = (1, 2, 3)$
 $B = (2, 2, 3)$ \rightarrow $A = (2, 2, 3)$
 $B = (2, 2, 3)$

Broadcast: Do **Problems 3** (5 mins)

Rule for figuring out behavior:

1. **Line up** array shapes starting from the **right**
2. **For each axis:**
 - a. If shapes match, continue to the left
 - b. If shapes don't match and one is 1, stretch its values to fit the larger
 - c. If shapes don't match and *neither* are 1, throw an error

```
np.arange(3) => [0, 1, 2]  
.reshape(3, 1) => [[0], [1], [2]]
```

Vectorization (4)

“Vectorization” means writing things with native NumPy operations rather than for loops

Much faster when possible!

Native C (low overhead) vs Python (high overhead)

SLOW

```
for i in range(H):  
    for j in range(W):  
        out[i,j,:] = (a[i,j] + b[i,j])/2.0
```

*Let me show you



Fast `(a + b) / 2.0`

Vectorization: do **Problems 4!**

“Vectorization” means writing things with native NumPy operations rather than for loops

```
np.mean( ... )
```

```
np.sum( ... )
```

Manipulating shapes (5)

Many times we want to shuffle the order of axes or combine them.

Remember arrays are *row-major*!

1	2	3
4	5	6
7	8	9

→

row-major								
1	2	3	4	5	6	7	8	9

Row-major order


a_{11}	a_{12}	a_{13}
a_{21}	a_{22}	a_{23}
a_{31}	a_{32}	a_{33}

Do problems 5.1 => 5.3

Thanks for coming!

Explore: Bonus & Einsum & Finish the rest.

Einsum



einsum examples!

```
>>> a = np.arange(4) # (4,)
```

```
array([0, 1, 4, 9]) # (4,)
```

```
>>> b = np.arange(4) # (4,)
```

```
>>> np.einsum('i,i->i', a, b)
```

```
array([[0, 0, 0, 0],
```

```
       [0, 1, 2, 3],
```

```
>>> np.einsum('i,j->ij', a, b)
```

```
       [0, 2, 4, 6],
```

```
>>> np.einsum('...i,...i->...', a, b)
```

```
       [0, 3, 6, 9]]) # (4, 4)
```

```
np.int64(14) # (,)
```